

My 4 Blocks: Voice & Chat Architecture

A guide to how typed chat and voice mode work—including system prompts, RAG retrieval, and information flow.

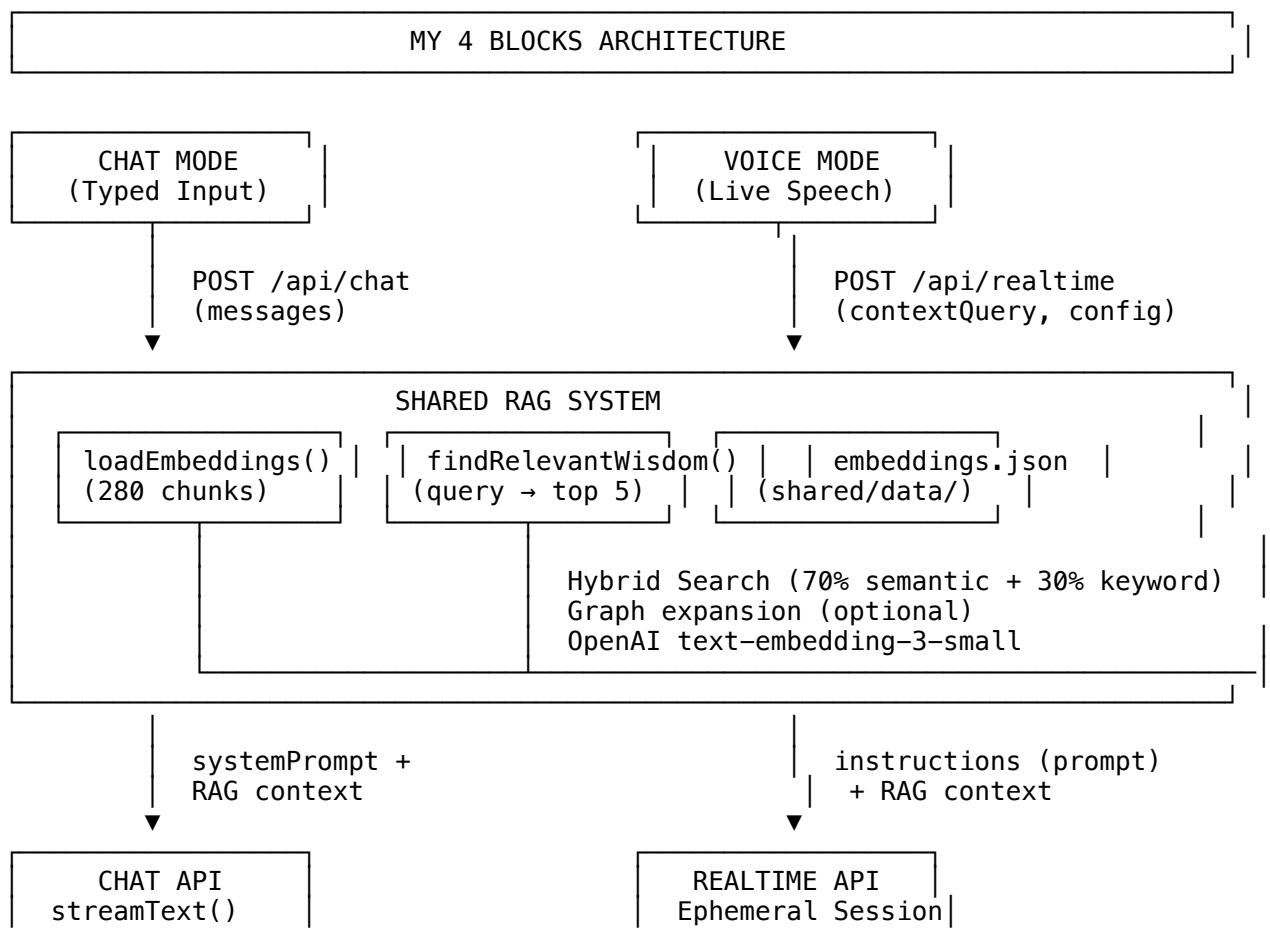
Overview

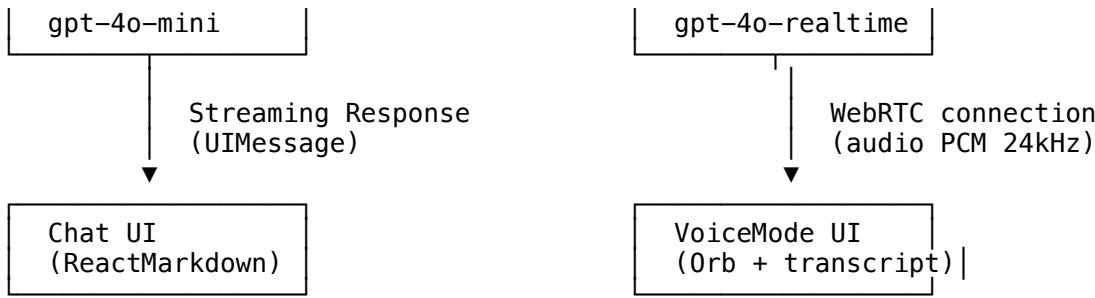
My 4 Blocks offers two ways to interact with the AI guide:

Mode	Input	Output	API
Chat	Typed text	Streamed text (markdown)	Vercel AI SDK → OpenAI Chat
Voice	Live speech	Real-time speech	OpenAI Realtime API (WebRTC)

Both modes share the same **RAG retrieval system** and **core knowledge base**, but differ in how they deliver that knowledge to the user.

Architecture Diagram





Chat Mode

Flow

1. **User types** → `sendMessage()` from `useChat()`
2. **Client** sends `POST /api/chat` with full message history
3. **Server (handleChatRequest):**
 - o Loads RAG embeddings (lazy init)
 - o Extracts query from last user message
 - o Calls `findRelevantWisdom(query, 5)` → hybrid search
 - o Builds `systemPrompt = SYSTEM_PROMPT + "\n\n## Relevant Book Context\n" + ragContext`
 - o Streams via `streamText()` with `system, messages, temperature, maxTokens`
4. **Response** streamed back as `UIMessage` stream
5. **Client** renders markdown with `ReactMarkdown`

System Prompt (Chat)

The chat system prompt is **richer** and includes:

- **Book structure** – Chapter outline
- **Four Blocks** – Nuanced definitions (Anger, Anxiety, Depression, Guilt)
- **Depression vs Guilt** – Exact distinction
- **ABC Model** – A–E (Activating Event, Belief, Consequence, Disputing, Effective new belief)
- **Seven Irrational Beliefs** – Full list
- **Three Insights** – Core principles
- **Narrator vs Observer** – Mindfulness framing
- **Communication style** – Warm, compassionate, non-judgmental
- **Key quotes** – Dr. Parr, Dōgen

RAG adds **top 5 relevant chunks** from the book under `## Relevant Book Context`.

Configuration

Option	Default	Description
<code>model</code>	<code>gpt-4o-mini</code>	Chat model
<code>temperature</code>	0.7	Response randomness
<code>maxTokens</code>	2000	Max tokens per response

ragEnabled	true	RAG retrieval
ragTopK	5	Chunks retrieved

Voice Mode

Flow

1. **User** starts voice session → **VoiceMode** component calls `POST /api/realtimetime`
2. **Request** body: `{ contextQuery?, config? }` (voice, style, ragEnabled, etc.)
3. **Server (handleRealtimeRequest):**
 - o Loads RAG embeddings (lazy init)
 - o Calls `buildVoiceInstructions(contextQuery, config):`
 - Builds voice system prompt with style
 - If `ragEnabled` and `contextQuery`, adds RAG context
 - o Creates ephemeral session via `POST https://api.openai.com/v1/realtimetime/sessions`
 - o Returns `{ client_secret, id, ... }` for WebRTC
4. **Client:**
 - o Uses `client_secret` to establish WebRTC connection
 - o Sends audio (PCM 24kHz) via WebRTC
 - o Receives audio stream, plays via `AudioContext`
 - o Shows transcript from data channel

System Prompt (Voice)

Voice uses the same knowledge base but **style-specific** instructions:

- **Style** – `direct` (default), `warm`, `casual`, `professional`
- **Book structure** – Same chapter outline
- **Four Blocks** – Same definitions
- **Depression vs Guilt** – Same distinction
- **ABC Model** – A → B → C
- **Seven Irrational Beliefs** – Short list
- **Three Insights** – Same core principles

RAG adds **top 5 relevant chunks** from the book under `## Relevant Book Context`.

Voice Style Presets

Style	Description
<code>direct</code>	Get to the point; no fluff; normal pace
<code>warm</code>	Friendly, supportive; natural rhythm
<code>casual</code>	Casual, conversational; everyday language
<code>professional</code>	Clear, structured; evidence-based

Voice Options

9 voices: `ash`, `alloy`, `ballad`, `coral`, `echo`, `marin`, `sage`, `shimmer`, `verse`.

Configuration

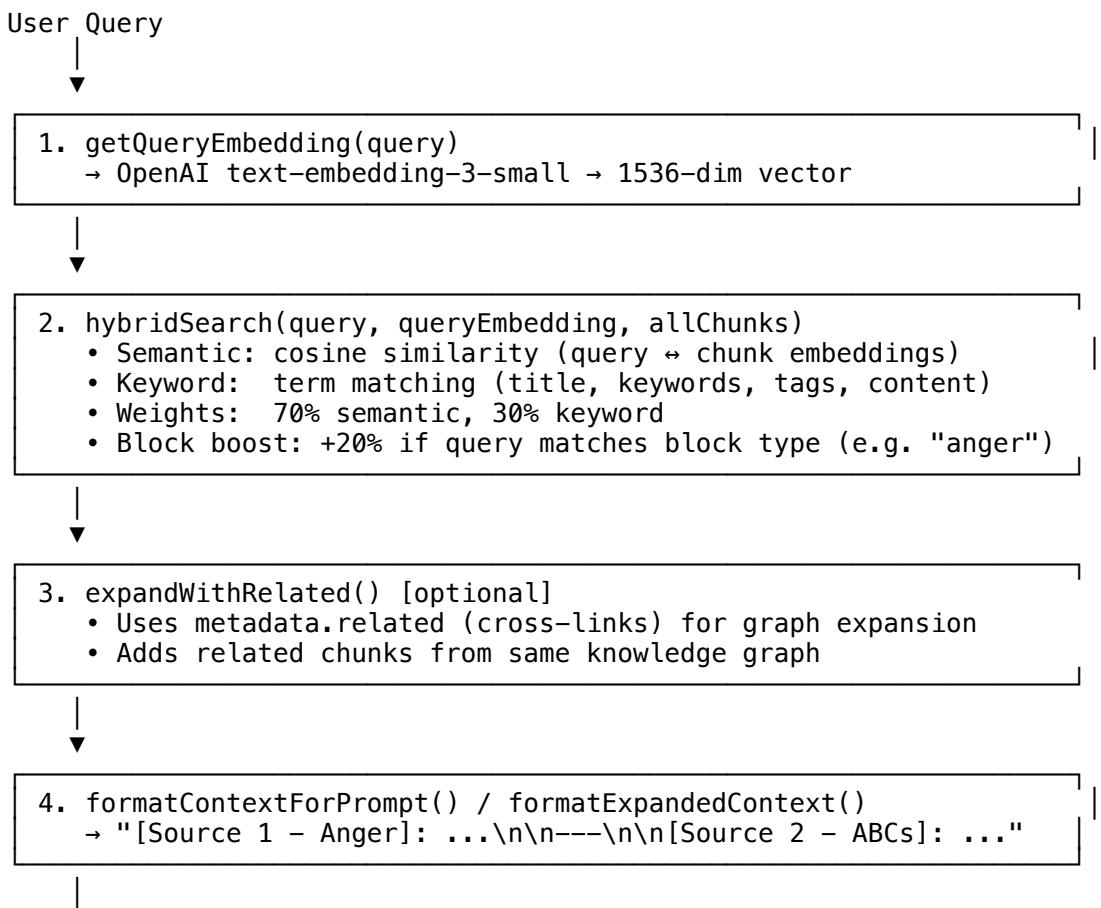
Option	Default	Description
voice	ash	TTS voice
style	direct	Conversation style
model	gpt-4o-realtime-preview-2024-12-17	Realtime model
ragEnabled	true	RAG retrieval
ragTopK	5	Chunks retrieved
turnDetection	semantic_vad, low	When to treat user as finished speaking

RAG: Retrieval-Augmented Generation

Data Source

- **Source:** full PDF of *You Only Have Four Problems* by Dr. Vincent E. Parr
- **Chunking:** Chonkie TokenChunker (500 tokens, 100 overlap)
- **Embeddings:** OpenAI text-embedding-3-small (1536 dims)
- **Storage:** shared/data/embeddings.json (~280 chunks)

Retrieval Pipeline



▼
Appended to system prompt as "## Relevant Book Context"

Fallback

If the embedding API fails, the system falls back to **keyword-only** search.

Information Management Summary

Aspect	Chat	Voice
Context window	Full message history sent each request	Ephemeral session; instructions fixed at session start
RAG trigger	Query = last user message	Query = optional contextQuery (e.g. suggested prompt)
Streaming	Text stream (UIMessage)	Audio stream (WebRTC)
Transcription	N/A	Whisper-1 on server
Turn detection	N/A	Semantic VAD (server decides when user finished)
Session persistence	Client maintains conversation	Session ends when WebRTC disconnected

File Reference

File	Purpose
shared/api/chat.ts	Chat API handler, SYSTEM_PROMPT, handleChatRequest
shared/api/realtime.ts	Voice API handler, buildSystemPrompt, createRealtimeSession
shared/lib/rag.ts	getRAGContext, findRelevantWisdom, loadEmbeddings
shared/lib/hybridSearch.ts	Hybrid semantic + keyword search
shared/lib/embeddings.ts	Query embedding (OpenAI or local)
shared/lib/vectorSearch.ts	Cosine similarity, formatContextForPrompt
shared/lib/keywordSearch.ts	Keyword search, emotion keyword expansion
shared/lib/graphExpansion.ts	Cross-link expansion via metadata.related
shared/data/embeddings.json	Pre-computed embeddings (280 chunks)
shared/components/VoiceMode.tsx	Voice UI, WebRTC, orb, transcript

Suggested Prompts (Chat Page)

When the user starts a chat, they can choose from suggested prompts:

- **Managing Anger** – "I keep getting angry at things I can't control. How can I stop?"
- **Understanding Anxiety** – "I'm worried about what might happen. How can I stop worrying?"

- **Overcoming Depression** – "I feel hopeless and unmotivated. What can I do?"
- **Releasing Guilt** – "I feel guilty about something I did. How can I let it go?"
- **The ABCs Model** – "How do my thoughts create my emotions?"
- **Core Beliefs** – "What are the irrational beliefs that cause suffering?"

These prompts are used both as **chat starters** and as **contextQuery** for voice sessions when the user starts voice mode from a suggested prompt.

Document generated for My 4 Blocks—the emotional wellness chat powered by the Four Blocks framework.